

# Lecture 1: Introduction

Statistical Learning and Data Mining

Xingye Qiao

Department of Mathematical Sciences  
Binghamton University

E-mail: [qiao@math.binghamton.edu](mailto:qiao@math.binghamton.edu)

Read: ISR Chs. 1–2 and ESLII Chs. 1–2

# Outline

- 1 What is statistical learning?
- 2 A framework for SL

The next section would be .....

1 What is statistical learning?

2 A framework for SL

# What is statistical learning?

- Intersects: Data science, statistics, applied math, computer science.
- How is statistical learning different from machine learning / artificial intelligence?
- Simple definition: ML = an algorithm that can learn from data without relying on rules-based programming;  
SL = formalization of relationships between variables in the form of mathematical equations.
- Different schools.
- Different eras.
- Assumptions involved.

## Name conventions

Names which refer to almost the same things

ML	Statistics
network, graphs	(generative) model
weight	parameter, coefficient
learning	fitting
generalization	test set performance
supervised learning	regression + classification
unsupervised learning	density estimator, clustering, and so on
learner, algorithm	model

# Statistical Learning Tasks

- Prediction.
- Exploratory Analysis (Data-Driven Discoveries).

# Types of statistical learning

- Supervised Learning: Data with labels or a response.
- Unsupervised Learning: Data with no labels.
- Others: Semi-supervised learning, recommender systems, online learning, network models, sequential learning, text mining, etc.

## Other categorizations

We have seen supervised vs unsupervised.

- Active vs. Passive learners. An active learner interacts with the environment at training time, while a passive learner only observes the information provided by the environment (or the teacher) without influencing or directing it.
- Online vs. Batch learning



## Course overview and expectations

- Piazza: homework assignments, lecture notes, discussions.
- Piazza activity counts.
- Deduction of 25% of the grade for homeworks that are not typeset using LaTeX. Deduction of 15% of the grade for each day homeworks are late (the final grade for a late homework that is  $N$  days late will be  $0.85^N$  times the real grade).
- Homeworks may be discussed with classmates but must be written and submitted individually.

The next section would be .....

1 What is statistical learning?

2 A framework for SL

## The SL framework

- Domain set:  $\mathcal{X}$ . Can be Euclidean space for simplicity.
- Label set:  $\mathcal{Y}$  can be a set of real numbers or a finite set
- Training data:  $S = \{(x_i, y_i)_{i=1, \dots, n} \in \mathcal{X} \times \mathcal{Y}\}$
- Unknown underlying distribution:  $(X, Y) \sim F$ .
- iid: Assume that each  $(X_i, Y_i)$  is iid according to  $F$ .
- Learner's output: a hypothesis or a prediction rule,  $h_n : \mathcal{X} \mapsto \mathcal{Y}$ . The subscript  $n$  emphasizes that the prediction rule depends on  $S$  (which has size  $n$ ) and hence is random.
- Measure of success: In regression,

$$R(h) = E_F(h(X) - Y)^2.$$

In classification,

$$R(h) = P_F(h(X) \neq Y).$$

Here  $h$  is general. May be deterministic, or may be data dependent  $h_n$  which is random. Synonymous names such as the generalization error, the test error, or the true error of  $h$ .

## Classification: Bayes rule

We now focus on binary classification ( $Y = \pm 1$ ) as an example.

- Given  $h$ , the risk of  $h$  is

$$R(h) = P_F(h(X) \neq Y) = E_F(\mathbb{1}\{h(X) \neq Y\}).$$

- Define the regression function

$$\eta(x) = E_F(Y|X = x) = 2P_F(Y = 1|X = x) - 1.$$

- Define a rule  $h_B(x) = \text{sign}(\eta(x))$
- $h_B$  achieves the minimal risk over all possible measurable functions:

$$R(h_B) = \inf_{h \in \mathcal{M}} R(h)$$

- We denote  $R(h_B)$  by  $R^*$ , called the **Bayes risk**.  $h_B$  is called the **Bayes rule**.
- Note  $h_B$  does not depend on the data, but on the unknown distribution  $F$ .

## Empirical Risk Minimization

Our goal is to identify  $h_B$ . However, the risk  $R(h)$  cannot be directly measured since  $F$  is unknown.

- It is common to find an approximation to the true mean, using the sample mean, such as approximating  $E_F(\mathbb{1}\{h(X) \neq Y\})$  by

$$R_n(h) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}\{h(X_i) \neq Y_i\}$$

- This is called the empirical risk, since this is equivalent to the risk with distribution  $F$  replaced by empirical distribution.
- The Empirical Risk Minimization (ERM) rule/algorithm finds a predictor  $h$  that minimizes  $R_n(h)$ . I.e.  $\operatorname{argmin}_h R_n(h)$
- Note that there may be several  $h$  that minimize the training error and an ERM algorithm is free to choose any such  $h$ .

# Overfitting

- Example:  $Y = \pm 1$ . Among 100 training data points uniformly distributed in the 2D sphere with radius 2,  $Y = 1$  deterministically as long as  $\|X\| < 1$ , or  $-1$  otherwise.
- We propose a possible ERM predictor (classifier):  $h(x) = y_i$  if there exists  $x_i = x$ , or 1 otherwise.
- Perfect for training data:  $R_n(h) = 0$ . Zero training error.
- Calculate the generalization error:  
 $R(h) = P_F(h(X) \neq Y) = 0.75$ . Large test error.
- This is overfitting: performance on the training set is excellent but performance on the true world is very bad.
- Intuitively, overfitting occurs when we can explain almost every data point in the training data.

## A possible solution

- Apply the ERM learning rule, but restrict the search space.
- Formally, the learner should choose in advance (before seeing the data) a set of predictors. This set is called a hypothesis class and is denoted by  $\mathcal{H}$ . That is, each  $h \in \mathcal{H}$  is a function mapping from  $\mathcal{X}$  to  $\mathcal{Y}$ . After deciding on  $\mathcal{H}$ , the learner samples a training set  $S$  and uses the ERM rule to choose a predictor out of the hypothesis class.
- It can be shown that this may reduce the generalization error.
- Why it might work: control the complexity of the model.




Figure: Trade-off between fit and complexity

Typically, one would look, in a collection of possible models, for one which fits well the data, but at the same time is as simple as possible



# Complexity

- There is no universal way of measuring simplicity (or its counterpart complexity) and the choice of a specific measure inherently depends on the problem at hand.
- It is actually in this choice that the designer of the learning algorithm introduces knowledge about the specific phenomenon under study.
- This leads to the statement that data cannot replace knowledge, or in pseudo-mathematical terms:

$$\text{Generalization} = \text{Data} + \text{Knowledge}$$

# Regularization

- Recall that the ERM seeks to find  $\operatorname{argmin}_h R_n(h)$
- Another approach: choosing a large model  $\mathcal{H}$  and then define on  $\mathcal{H}$  a regularizer, typically a norm  $\|h\|$ . Then one has to minimize the **regularized** empirical risk:

$$h_n = \operatorname{argmin}_{h \in \mathcal{H}} R_n(h) + \lambda \|h\|^2$$

- There is free parameter  $\lambda$ , called the **regularization parameter** which allows to choose the right trade-off between fit and complexity.
- Tuning  $\lambda$  is usually a hard problem and most often, one uses extra validation data for this task.
- Most existing (and successful) methods can be thought of as regularization methods.

## A Simple Approach to Prediction: Nearest Neighbors

The  $k$ -nearest neighbor classifiers are memory based, and require no model to fit. Given a point  $\mathbf{x}^*$ , we find  $k$  points  $\mathbf{x}_{(r)}, r = 1, \dots, k$ , among training inputs, closest in *distance* to  $\mathbf{x}^*$ .  $\mathbf{x}^*$  is classified using majority vote among the  $k$  neighbors.

- simple to use.
- shown to be successful in real examples.
- requires large memory if the dataset is huge.
- for regression tasks, take the average of the  $y$  values of the  $k$  neighbors to predict  $y^*$

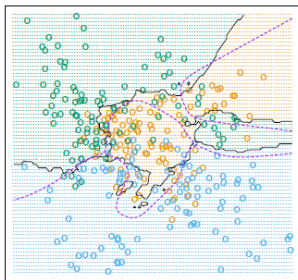
Example from ESLII (Hastie, Ribshirani, Friedman)

$k$ -Nearest Neighbor classifiers applied to a simulated data set, with three groups.


The decision boundary of a 15-Nearest Neighbor classifier (top) is fairly smooth compared to a 1-Nearest Neighbor classifier (bottom).

Smaller  $k$ , more complex model.  
Larger  $k$ , more simple model.

15-Nearest Neighbors



1-Nearest Neighbor



## lecture1 R code

- 1 Predict the label based on the labels of the  $K$  closest observations to the data point of interest.
- 2 What happens to the test prediction error when  $K$  is small?  
Large?

## High-dimensional data & curse of dimensionality

- Intuition breaks down in high dimensions, and the phenomenon is commonly referred to as the curse of dimensionality
- Intuition 1: a small local neighborhood with  $s$  fraction of the entire space should approximate the true value well. However it has a radius of  $s^{1/p}$  which goes to 1 quickly as  $p$  increases, hence **no longer local**.
- Intuition 2: it is easy to reach a neighbor. However, for large  $p$ , most data points are closer to the edge of the sample space than to any neighbor. (Not a good thing.)
- Local methods such as  $k$ NN can fail for high dimensional data.

# Function Approximation

- Many SL problems can be seen as function approximation.
- $(x_i, y_i)$  is seen as a point in a  $(p + 1)$ -dimensional space. The function to be approximated  $f(x)$  is related to the data via a model  $y_i = f(x_i) + \varepsilon_i$
- While this looks familiar from regression, it is also related to classification problems via  $y_i = \text{sign}(f(x_i) + \varepsilon_i)$
- Common class of approximations:
  - Linear model:  $f(x) = x^T \beta$
  - Linear basis expansion:  $f(x) = \sum_{k=1}^K h_k(x) \theta_k$  where  $h_k$  are a suitable set of functions or transformations of the input vector  $x$  such as  $x_1^2$ ,  $x_1 x_2^2$ ,  $\cos(x_1)$ , etc.
- To this end, finding the optimal  $h$  boils down to  $f$ , then to  $\beta$  (or  $\theta_k$ 's).

## Summary

- What is SL?
- Framework of SL. Training data. Goal. Risk function.
- Bayes rule. Training and test errors.
- Overfitting
- Model complexity and regularization.
- Curse of dimensionality
- $k$ NN – an example.